

# Formatting Instructions For NeurIPS 2025

Anonymous Author(s)

### Affiliation, Address

anon.email@example.org

### Abstract

The abstract paragraph should be indented  $\frac{1}{2}$  inch (3 picas) on both the left- and right-hand margins. Use 10 point type, with a vertical spacing (leading) of 11 points. The word **Abstract** must be centered, bold, and in point size 12. Two line spaces precede the abstract. The abstract must be limited to one paragraph.

## **5 1 Submission of papers to NeurIPS 2025**

6 Please read the instructions below carefully and follow them faithfully.

7 1.1 Style

8 Papers to be submitted to NeurIPS 2025 must be prepared according to the instructions presented  
9 here. Papers may only be up to **nine** pages long, including figures. Additional pages *containing only*  
10 *acknowledgments and references* are allowed. Papers that exceed the page limit will not be reviewed,  
11 or in any other way considered for presentation at the conference.

12 The margins in 2025 are the same as those in previous years.

<sup>13</sup> Authors are required to use the NeurIPS L<sup>A</sup>T<sub>E</sub>X style files obtainable at the NeurIPS website as  
<sup>14</sup> indicated below. Please make sure you use the current files and not previous versions. Tweaking the  
<sup>15</sup> style files may be grounds for rejection.

16 1.2 Retrieval of style files

<sup>17</sup> The style files for NeurIPS and other conference information are available on the website at [www.neurips.cc](http://www.neurips.cc)

18 <http://www.neurips.cc/>

<sup>19</sup> The file `neurips_2025.pdf` contains these instructions and illustrates the various formatting requirements your NeurIPS paper must satisfy.  
<sup>20</sup>

21 The only supported style file for NeurIPS 2025 is `neurips_2025.sty`, rewritten for L<sup>A</sup>T<sub>E</sub>X 2 <sub>$\epsilon$</sub> .  
22 **Previous style files for L<sup>A</sup>T<sub>E</sub>X 2.09, Microsoft Word, and RTF are no longer supported!**

<sup>23</sup> The L<sup>A</sup>T<sub>E</sub>X style file contains three optional arguments: `final`, which creates a camera-ready copy,  
<sup>24</sup> `preprint`, which creates a preprint for submission to, e.g., arXiv, and `nonatbib`, which will not load  
<sup>25</sup> the `natbib` package for you in case of package clash.

**Preprint option** If you wish to post a preprint of your work online, e.g., on arXiv, using the NeurIPS style, please use the `preprint` option. This will create a nonanonymized version of your work with the text “Preprint. Work in progress.” in the footer. This version may be distributed as you see fit, as long as you do not say which conference it was submitted to. Please **do not** use the `final` option, which should **only** be used for papers accepted to NeurIPS.

- 31 At submission time, please omit the `final` and `preprint` options. This will anonymize your  
32 submission and add line numbers to aid review. Please do *not* refer to these line numbers in your  
33 paper as they will be removed during generation of camera-ready copies.
- 34 The file `neurips_2025.tex` may be used as a “shell” for writing your paper. All you have to do is  
35 replace the author, title, abstract, and text of the paper with your own.
- 36 The formatting instructions contained in these style files are summarized in Sections 2, 3, and 4  
37 below.

## 38 **2 General formatting instructions**

- 39 The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long.  
40 The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing (leading) of 11 points.  
41 Times New Roman is the preferred typeface throughout, and will be selected for you by default.  
42 Paragraphs are separated by  $\frac{1}{2}$  line space (5.5 points), with no indentation.
- 43 The paper title should be 17 point, initial caps/lower case, bold, centered between two horizontal  
44 rules. The top rule should be 4 points thick and the bottom rule should be 1 point thick. Allow  $\frac{1}{4}$  inch  
45 space above and below the title to rules. All pages should start at 1 inch (6 picas) from the top of  
46 the page.
- 47 For the final version, authors’ names are set in boldface, and each name is centered above the  
48 corresponding address. The lead author’s name is to be listed first (left-most), and the co-authors’  
49 names (if different address) are set to follow. If there is only one co-author, list both author and co-  
50 author side by side.
- 51 Please pay special attention to the instructions in Section 4 regarding figures, tables, acknowledg-  
52 ments, and references.

## 53 **3 Headings: first level**

- 54 All headings should be lower case (except for first word and proper nouns), flush left, and bold.  
55 First-level headings should be in 12-point type.

### 56 **3.1 Headings: second level**

- 57 Second-level headings should be in 10-point type.

#### 58 **3.1.1 Headings: third level**

- 59 Third-level headings should be in 10-point type.

60 **Paragraphs** There is also a `\paragraph` command available, which sets the heading in bold, flush  
61 left, and inline with the text, with the heading followed by `1em` of space.

## 62 **4 Citations, figures, tables, references**

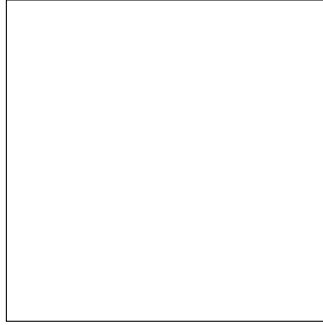
- 63 These instructions apply to everyone.

### 64 **4.1 Citations within the text**

65 The `natbib` package will be loaded for you by default. Citations may be author/year or numeric, as  
66 long as you maintain internal consistency. As to the format of the references themselves, any style  
67 is acceptable as long as it is used consistently.

- 68 The documentation for `natbib` may be found at

69 <http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf>



70

Figure 1: Sample figure caption.

71 Of note is the command `\citet`, which produces citations appropriate for use in inline text. For  
72 example,

73 `\citet{hasselmo} investigated\ldots`

74 produces

75 Hasselmo, et al. (1995) investigated\ldots

76 If you wish to load the `natbib` package with options, you may add the following before loading the  
77 `neurips_2025` package:

78 `\PassOptionsToPackage{options}{natbib}`

79 If `natbib` clashes with another package you load, you can add the optional argument `nonatbib` when  
80 loading the style file:

81 `\usepackage[nonatbib]{neurips_2025}`

82 As submission is double blind, refer to your own published work in the third person. That is, use “In  
83 the previous work of Jones et al. [4],” not “In our previous work [4].” If you cite your other papers  
84 that are not widely available (e.g., a journal paper under review), use anonymous author names in the  
85 citation, e.g., an author of the form “A. Anonymous” and include a copy of the anonymized paper  
86 in the supplementary material.

87 **4.2 Footnotes**

88 Footnotes should be used sparingly. If you do require a footnote, indicate footnotes with a number<sup>1</sup>  
89 in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote  
90 with a horizontal rule of 2 inches (12 picas).

91 Note that footnotes are properly typeset *after* punctuation marks.<sup>2</sup>

92 **4.3 Figures**

93 All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of repro-  
94 duction. The figure number and caption always appear after the figure. Place one line space before  
95 the figure caption and one line space after the figure. The figure caption should be lower case (except  
96 for first word and proper nouns); figures are numbered consecutively.

97 You may use color figures. However, it is best for the figure captions and the paper body to be legible  
98 if the paper is printed in either black/white or in color.

---

<sup>1</sup>Sample of the first footnote.

<sup>2</sup>As in this example.

99

Table 1: Sample table title.

---

Part			
	Name	Description	Size ( $\mu\text{m}$ )
102	Dendrite	Input terminal	$\sim 100$
103	Axon	Output terminal	$\sim 10$
104	Soma	Cell body	up to $10^6$

---

105 **4.4 Tables**106 All tables must be centered, neat, clean and legible. The table number and title always appear before  
107 the table. See Table 1.108 Place one line space before the table title, one line space after the table title, and one line space  
109 after the table. The table title must be lower case (except for first word and proper nouns); tables are  
110 numbered consecutively.111 Note that publication-quality tables *do not contain vertical rules*. We strongly suggest the use of the  
112 `booktabs` package, which allows for typesetting high-quality, professional tables:113 <https://www.ctan.org/pkg/booktabs>

114 This package was used to typeset Table 1.

115 **4.5 Math**116 Note that display math in bare TeX commands will not create correct line numbers  
117 for submission. Please use LaTeX (or AMSTeX) commands for unnumbered display  
118 math. (You really shouldn't be using \$\$ anyway; see <https://tex.stackexchange.com/questions/503/why-is-preferable-to> and <https://tex.stackexchange.com/questions/40492/what-are-the-differences-between-align-equation-and-displaymath> for more information.)122 **4.6 Final instructions**123 Do not change any aspects of the formatting parameters in the style files. In particular, do not modify  
124 the width or length of the rectangle the text should fit into, and do not change font sizes (except  
125 perhaps in the **References** section; see below). Please note that pages should be numbered.126 **5 Preparing PDF files**

127 Please prepare submission files with paper size “US Letter,” and not, for example, “A4.”

128 Fonts were the main cause of problems in the past years. Your PDF file must only contain Type 1 or  
129 Embedded TrueType fonts. Here are a few instructions to achieve this.

- 130 • You should directly generate PDF files using `pdflatex`.
- 131 • You can check which fonts a PDF file uses. In Acrobat Reader, select the menu Files  
132 > Document Properties > Fonts and select Show All Fonts. You can also use the program  
133 `pdffonts` which comes with `xpdf` and is available out-of-the-box on most Linux machines.
- 134 • `xfig` “patterned” shapes are implemented with bitmap fonts. Use “solid” shapes instead.
- 135 • The `\bbold` package almost always uses bitmap fonts. You should use the equivalent AMS  
136 Fonts:

137 `\usepackage{amsfonts}`

138 followed by, e.g., `\mathbb{R}`, `\mathbb{N}`, or `\mathbb{C}` for  $\mathbb{R}$ ,  $\mathbb{N}$  or  $\mathbb{C}$ . You can also  
139 use the following workaround for reals, natural and complex:

140       `\newcommand{\RR}{\mathbb{R}} %real numbers`  
141       `\newcommand{\Nat}{\mathbb{N}} %natural numbers`  
142       `\newcommand{\CC}{\mathbb{C}} %complex numbers`

143 Note that `amsfonts` is automatically loaded by the `amssymb` package.

144 If your file contains Type 3 fonts or non embedded TrueType fonts, we will ask you to fix it.

## 145 **5.1 Margins in L<sup>A</sup>T<sub>E</sub>X**

146 Most of the margin problems come from figures positioned by hand using `\special` or other  
147 commands. We suggest using the command `\includegraphics` from the `graphicx` package. Always  
148 specify the figure width as a multiple of the line width as in the example below:

149       `\usepackage[pdf]{graphicx} ...`  
150       `\includegraphics[width=0.8\linewidth]{myfile.pdf}`

151 See Section 4.4 in the graphics bundle documentation (<http://mirrors.ctan.org/macros/latex/required/graphics/grfguide.pdf>)

153 A number of width problems arise when L<sup>A</sup>T<sub>E</sub>X cannot properly hyphenate a line. please give L<sup>A</sup>T<sub>E</sub>X  
154 hyphenation hints using the `\-` command when necessary.

## 155 **References**

156 References follow the acknowledgments in the camera-ready paper. Use unnumbered first-level  
157 heading for the references. Any choice of citation style is acceptable as long as you are consistent.  
158 It is permissible to reduce the font size to `small` (9 point) when listing the references. Note that the  
159 Reference section does not count towards the page limit.

- 160 [1] J. Alexander and M. C. Mozer, “Template-based algorithms for connectionist rule extraction,” *Advances*  
161     in *neural information processing systems*, vol. 7, 1994.
- 162 [2] J. M. Bower and D. Beeman, *The book of GENESIS: exploring realistic neural models with the GEneral*  
163     *NEural SImulation System*. Springer Science & Business Media, 2012.
- 164 [3] M. E. Hasselmo, E. Schnell, and E. Barkai, “Dynamics of learning and recall at excitatory recurrent  
165     synapses and cholinergic modulation in rat hippocampal region CA3,” *Journal of Neuroscience*, vol. 15,  
166     no. 7, pp. 5249–5262, 1995.

## 167 **A Appendix / supplemental material**

168 Optionally include supplemental material (complete proofs, additional experiments and plots) in  
169 appendix. All such materials **SHOULD be included in the main submission**.

170 **NeurIPS Paper Checklist**

171 The checklist is designed to encourage best practices for responsible machine learning research,  
172 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not  
173 remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist  
174 should follow the references and precede the (optional) supplemental material. The checklist does  
175 NOT count towards the page limit.

176 Please read the checklist guidelines carefully for information on how to answer these questions. For  
177 each question in the checklist:

- 178 • You should answer [YES], [NO], or [NA].
- 179 • [NA] means either that the question is Not Applicable for that particular paper or the  
180 relevant information is Not Available.
- 181 • Please provide a short (1–2 sentence) justification right after your answer (even for NA).

182 **The checklist answers are an integral part of your paper submission.** They are visible to the  
183 reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it  
184 (after eventual revisions) with the final version of your paper, and its final version will be published  
185 with the paper.

186 The reviewers of your paper will be asked to use the checklist as one of the factors in their evalua-  
187 tion. While “[YES]” is generally preferable to “[NO]”, it is perfectly acceptable to answer “[NO]”  
188 provided a proper justification is given (e.g., “error bars are not reported because it would be too  
189 computationally expensive” or “we were unable to find the license for the dataset we used”). In  
190 general, answering “[NO]” or “[NA]” is not grounds for rejection. While the questions are phrased  
191 in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your  
192 best judgment and write a justification to elaborate. All supporting evidence can appear either in the  
193 main paper or the supplemental material, provided in appendix. If you answer [YES] to a question,  
194 in the justification please point to the section(s) where related material for the question can be found.

195 **IMPORTANT**, please:

- 196 • **Delete this instruction block, but keep the section heading “NeurIPS paper checklist”,**
- 197 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 198 • **Do not modify the questions and only use the provided macros for your answers.**

199 **1. Claims**

200 Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s  
201 contributions and scope?

202 Answer: [TODO]

203 Justification: [TODO]

204 Guidelines:

- 205 • The answer NA means that the abstract and introduction do not include the claims made in  
206 the paper.
- 207 • The abstract and/or introduction should clearly state the claims made, including the contri-  
208 butions made in the paper and important assumptions and limitations. A No or NA answer  
209 to this question will not be perceived well by the reviewers.
- 210 • The claims made should match theoretical and experimental results, and reflect how much  
211 the results can be expected to generalize to other settings.
- 212 • It is fine to include aspirational goals as motivation as long as it is clear that these goals are  
213 not attained by the paper.

214 **2. Limitations**

215      Question: Does the paper discuss the limitations of the work performed by the authors?

216      Answer: [TODO]

217      Justification: [TODO]

218      Guidelines:

- 219      • The answer NA means that the paper has no limitation while the answer No means that the  
220      paper has limitations, but those are not discussed in the paper.
- 221      • The authors are encouraged to create a separate “Limitations” section in their paper.
- 222      • The paper should point out any strong assumptions and how robust the results are to  
223      violations of these assumptions (e.g., independence assumptions, noiseless settings, model  
224      well-specification, asymptotic approximations only holding locally). The authors should  
225      reflect on how these assumptions might be violated in practice and what the implications  
226      would be.
- 227      • The authors should reflect on the scope of the claims made, e.g., if the approach was only  
228      tested on a few datasets or with a few runs. In general, empirical results often depend on  
229      implicit assumptions, which should be articulated.
- 230      • The authors should reflect on the factors that influence the performance of the approach.  
231      For example, a facial recognition algorithm may perform poorly when image resolution  
232      is low or images are taken in low lighting. Or a speech-to-text system might not be used  
233      reliably to provide closed captions for online lectures because it fails to handle technical  
234      jargon.
- 235      • The authors should discuss the computational efficiency of the proposed algorithms and  
236      how they scale with dataset size.
- 237      • If applicable, the authors should discuss possible limitations of their approach to address  
238      problems of privacy and fairness.
- 239      • While the authors might fear that complete honesty about limitations might be used by  
240      reviewers as grounds for rejection, a worse outcome might be that reviewers discover  
241      limitations that aren’t acknowledged in the paper. The authors should use their best  
242      judgment and recognize that individual actions in favor of transparency play an important  
243      role in developing norms that preserve the integrity of the community. Reviewers will be  
244      specifically instructed to not penalize honesty concerning limitations.

245      **3. Theory Assumptions and Proofs**

246      Question: For each theoretical result, does the paper provide the full set of assumptions and a  
247      complete (and correct) proof?

248      Answer: [TODO]

249      Justification: [TODO]

250      Guidelines:

- 251      • The answer NA means that the paper does not include theoretical results.
- 252      • All the theorems, formulas, and proofs in the paper should be numbered and cross-  
253      referenced.
- 254      • All assumptions should be clearly stated or referenced in the statement of any theorems.
- 255      • The proofs can either appear in the main paper or the supplemental material, but if they  
256      appear in the supplemental material, the authors are encouraged to provide a short proof  
257      sketch to provide intuition.
- 258      • Inversely, any informal proof provided in the core of the paper should be complemented by  
259      formal proofs provided in appendix or supplemental material.
- 260      • Theorems and Lemmas that the proof relies upon should be properly referenced.

261      **4. Experimental Result Reproducibility**

262 Question: Does the paper fully disclose all the information needed to reproduce the main  
263 experimental results of the paper to the extent that it affects the main claims and/or conclusions  
264 of the paper (regardless of whether the code and data are provided or not)?

265 Answer: [TODO]

266 Justification: [TODO]

267 Guidelines:

- 268 • The answer NA means that the paper does not include experiments.
- 269 • If the paper includes experiments, a No answer to this question will not be perceived well  
270 by the reviewers: Making the paper reproducible is important, regardless of whether the  
271 code and data are provided or not.
- 272 • If the contribution is a dataset and/or model, the authors should describe the steps taken to  
273 make their results reproducible or verifiable.
- 274 • Depending on the contribution, reproducibility can be accomplished in various ways. For  
275 example, if the contribution is a novel architecture, describing the architecture fully might  
276 suffice, or if the contribution is a specific model and empirical evaluation, it may be  
277 necessary to either make it possible for others to replicate the model with the same dataset,  
278 or provide access to the model. In general, releasing code and data is often one good way  
279 to accomplish this, but reproducibility can also be provided via detailed instructions for  
280 how to replicate the results, access to a hosted model (e.g., in the case of a large language  
281 model), releasing of a model checkpoint, or other means that are appropriate to the research  
282 performed.
- 283 • While NeurIPS does not require releasing code, the conference does require all submissions  
284 to provide some reasonable avenue for reproducibility, which may depend on the nature of  
285 the contribution. For example
  - 286 (a) If the contribution is primarily a new algorithm, the paper should make it clear how to  
287 reproduce that algorithm.
  - 288 (b) If the contribution is primarily a new model architecture, the paper should describe the  
289 architecture clearly and fully.
  - 290 (c) If the contribution is a new model (e.g., a large language model), then there should  
291 either be a way to access this model for reproducing the results or a way to reproduce  
292 the model (e.g., with an open-source dataset or instructions for how to construct the  
293 dataset).
  - 294 (d) We recognize that reproducibility may be tricky in some cases, in which case authors  
295 are welcome to describe the particular way they provide for reproducibility. In the case  
296 of closed-source models, it may be that access to the model is limited in some way  
297 (e.g., to registered users), but it should be possible for other researchers to have some  
298 path to reproducing or verifying the results.

299 **5. Open access to data and code**

300 Question: Does the paper provide open access to the data and code, with sufficient instructions  
301 to faithfully reproduce the main experimental results, as described in supplemental material?

302 Answer: [TODO]

303 Justification: [TODO]

304 Guidelines:

- 305 • The answer NA means that paper does not include experiments requiring code.
- 306 • Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 307 • While we encourage the release of code and data, we understand that this might not be  
308 possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not includ-  
309 ing code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- 311     • The instructions should contain the exact command and environment needed to run to  
312       reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.  
313  
314     • The authors should provide instructions on data access and preparation, including how to  
315       access the raw data, preprocessed data, intermediate data, and generated data, etc.  
316  
317     • The authors should provide scripts to reproduce all experimental results for the new pro-  
318       posed method and baselines. If only a subset of experiments are reproducible, they should  
319       state which ones are omitted from the script and why.  
320  
321     • At submission time, to preserve anonymity, the authors should release anonymized versions  
322       (if applicable).  
323     • Providing as much information as possible in supplemental material (appended to the paper)  
324       is recommended, but including URLs to data and code is permitted.

325     **6. Experimental Setting/Details**

326     Question: Does the paper specify all the training and test details (e.g., data splits, hyperpara-  
327       meters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

328     Answer: [TODO]

329     Justification: [TODO]

330     Guidelines:

- 331     • The answer NA means that the paper does not include experiments.  
332     • The experimental setting should be presented in the core of the paper to a level of detail  
333       that is necessary to appreciate the results and make sense of them.  
334     • The full details can be provided either with the code, in appendix, or as supplemental  
335       material.

336     **7. Experiment Statistical Significance**

337     Question: Does the paper report error bars suitably and correctly defined or other appropriate  
338       information about the statistical significance of the experiments?

339     Answer: [TODO]

340     Justification: [TODO]

341     Guidelines:

- 342     • The answer NA means that the paper does not include experiments.  
343     • The authors should answer “Yes” if the results are accompanied by error bars, confidence  
344       intervals, or statistical significance tests, at least for the experiments that support the main  
345       claims of the paper.  
346     • The factors of variability that the error bars are capturing should be clearly stated (for  
347       example, train/test split, initialization, random drawing of some parameter, or overall run  
348       with given experimental conditions).  
349     • The method for calculating the error bars should be explained (closed form formula, call to  
350       a library function, bootstrap, etc.)  
351     • The assumptions made should be given (e.g., Normally distributed errors).  
352     • It should be clear whether the error bar is the standard deviation or the standard error of  
353       the mean.  
354     • It is OK to report 1-sigma error bars, but one should state it. The authors should preferably  
355       report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality  
356       of errors is not verified.  
357     • For asymmetric distributions, the authors should be careful not to show in tables or figures  
358       symmetric error bars that would yield results that are out of range (e.g. negative error rates).

- 357       • If error bars are reported in tables or plots, The authors should explain in the text how they  
358        were calculated and reference the corresponding figures or tables in the text.

359   **8. Experiments Compute Resources**

360   Question: For each experiment, does the paper provide sufficient information on the computer  
361    resources (type of compute workers, memory, time of execution) needed to reproduce the  
362    experiments?

363   Answer: [TODO]

364   Justification: [TODO]

365   Guidelines:

- 366       • The answer NA means that the paper does not include experiments.  
367       • The paper should indicate the type of compute workers CPU or GPU, internal cluster, or  
368        cloud provider, including relevant memory and storage.  
369       • The paper should provide the amount of compute required for each of the individual  
370        experimental runs as well as estimate the total compute.  
371       • The paper should disclose whether the full research project required more compute than the  
372        experiments reported in the paper (e.g., preliminary or failed experiments that didn't make  
373        it into the paper).

374   **9. Code Of Ethics**

375   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS  
376    Code of Ethics <https://neurips.cc/public/EthicsGuidelines>

377   Answer: [TODO]

378   Justification: [TODO]

379   Guidelines:

- 380       • The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.  
381       • If the authors answer No, they should explain the special circumstances that require a  
382        deviation from the Code of Ethics.  
383       • The authors should make sure to preserve anonymity (e.g., if there is a special consideration  
384        due to laws or regulations in their jurisdiction).

385   **10. Broader Impacts**

386   Question: Does the paper discuss both potential positive societal impacts and negative societal  
387    impacts of the work performed?

388   Answer: [TODO]

389   Justification: [TODO]

390   Guidelines:

- 391       • The answer NA means that there is no societal impact of the work performed.  
392       • If the authors answer NA or No, they should explain why their work has no societal impact  
393        or why the paper does not address societal impact.  
394       • Examples of negative societal impacts include potential malicious or unintended uses  
395        (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g.,  
396        deployment of technologies that could make decisions that unfairly impact specific groups),  
397        privacy considerations, and security considerations.  
398       • The conference expects that many papers will be foundational research and not tied to  
399        particular applications, let alone deployments. However, if there is a direct path to any  
400        negative applications, the authors should point it out. For example, it is legitimate to point

401 out that an improvement in the quality of generative models could be used to generate  
402 deepfakes for disinformation. On the other hand, it is not needed to point out that a generic  
403 algorithm for optimizing neural networks could enable people to train models that generate  
404 Deepfakes faster.

- 405 • The authors should consider possible harms that could arise when the technology is being  
406 used as intended and functioning correctly, harms that could arise when the technology is  
407 being used as intended but gives incorrect results, and harms following from (intentional  
408 or unintentional) misuse of the technology.
- 409 • If there are negative societal impacts, the authors could also discuss possible mitigation  
410 strategies (e.g., gated release of models, providing defenses in addition to attacks, mecha-  
411 nisms for monitoring misuse, mechanisms to monitor how a system learns from feedback  
412 over time, improving the efficiency and accessibility of ML).

413 **11. Safeguards**

414 Question: Does the paper describe safeguards that have been put in place for responsible release  
415 of data or models that have a high risk for misuse (e.g., pretrained language models, image  
416 generators, or scraped datasets)?

417 Answer: [TODO]

418 Justification: [TODO]

419 Guidelines:

- 420 • The answer NA means that the paper poses no such risks.
- 421 • Released models that have a high risk for misuse or dual-use should be released with  
422 necessary safeguards to allow for controlled use of the model, for example by requiring  
423 that users adhere to usage guidelines or restrictions to access the model or implementing  
424 safety filters.
- 425 • Datasets that have been scraped from the Internet could pose safety risks. The authors  
426 should describe how they avoided releasing unsafe images.
- 427 • We recognize that providing effective safeguards is challenging, and many papers do not  
428 require this, but we encourage authors to take this into account and make a best faith effort.

429 **12. Licenses for existing assets**

430 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the  
431 paper, properly credited and are the license and terms of use explicitly mentioned and properly  
432 respected?

433 Answer: [TODO]

434 Justification: [TODO]

435 Guidelines:

- 436 • The answer NA means that the paper does not use existing assets.
- 437 • The authors should cite the original paper that produced the code package or dataset.
- 438 • The authors should state which version of the asset is used and, if possible, include a URL.
- 439 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 440 • For scraped data from a particular source (e.g., website), the copyright and terms of service  
441 of that source should be provided.
- 442 • If assets are released, the license, copyright information, and terms of use in the package  
443 should be provided. For popular datasets, <https://paperswithcode.com/datasets> has  
444 curated licenses for some datasets. Their licensing guide can help determine the license of  
445 a dataset.
- 446 • For existing datasets that are re-packaged, both the original license and the license of the  
447 derived asset (if it has changed) should be provided.

- 448       • If this information is not available online, the authors are encouraged to reach out to the  
449       asset's creators.

450     **13. New Assets**

451     Question: Are new assets introduced in the paper well documented and is the documentation  
452     provided alongside the assets?

453     Answer: [TODO]

454     Justification: [TODO]

455     Guidelines:

- 456       • The answer NA means that the paper does not release new assets.  
457       • Researchers should communicate the details of the dataset/code/model as part of their  
458       submissions via structured templates. This includes details about training, license, limita-  
459       tions, etc.  
460       • The paper should discuss whether and how consent was obtained from people whose asset  
461       is used.  
462       • At submission time, remember to anonymize your assets (if applicable). You can either  
463       create an anonymized URL or include an anonymized zip file.

464     **14. Crowdsourcing and Research with Human Subjects**

465     Question: For crowdsourcing experiments and research with human subjects, does the paper  
466     include the full text of instructions given to participants and screenshots, if applicable, as well  
467     as details about compensation (if any)?

468     Answer: [TODO]

469     Justification: [TODO]

470     Guidelines:

- 471       • The answer NA means that the paper does not involve crowdsourcing nor research with  
472       human subjects.  
473       • Including this information in the supplemental material is fine, but if the main contribution  
474       of the paper involves human subjects, then as much detail as possible should be included  
475       in the main paper.  
476       • According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or  
477       other labor should be paid at least the minimum wage in the country of the data collector.

478     **15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human  
479       Subjects**

480     Question: Does the paper describe potential risks incurred by study participants, whether such  
481     risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals  
482     (or an equivalent approval/review based on the requirements of your country or institution) were  
483     obtained?

484     Answer: [TODO]

485     Justification: [TODO]

486     Guidelines:

- 487       • The answer NA means that the paper does not involve crowdsourcing nor research with  
488       human subjects.  
489       • Depending on the country in which research is conducted, IRB approval (or equivalent)  
490       may be required for any human subjects research. If you obtained IRB approval, you should  
491       clearly state this in the paper.

- 492 • We recognize that the procedures for this may vary significantly between institutions and  
493 locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines  
494 for their institution.
- 495 • For initial submissions, do not include any information that would break anonymity (if  
496 applicable), such as the institution conducting the review.